




Embracing the log in UI engineering

July 3, 2015

This is drawn from a part of my talk at React Europe on July 3rd, 2015. It's about why using a log to record data within a UI can enable some awesome product features.



Hi! My name's Kevin (@krob), and today I'm going to talk about why it can be useful to use a log as the central concept for organizing data and computation when building user interfaces.

Inspiration from the backend



First, let's talk about what this looks like in backend systems, to set some context for what I'm talking about.



On the team I work on, Answers, we use immutable logs in several places in our backend systems. This has a lot of benefits, and I think some of these apply directly to UI engineering even though they are quite different problem spaces.

<https://blog.twitter.com/2015/handling-five-billion-sessions-a-day-in-real-time>

Embrace the log





Diagram adapted from <http://blog.confluent.io/2015/03/04/>



As writes come in, they're written to a log. There's no computation or mutation.

Embrace the log



Then, separately, processing components can read from the log and compute whatever "views" they need of the data. These might be filters, aggregations, or special indexes.

Embrace the log

Writes:


~~mutate state in place~~

append-only stream of
immutable facts




A central idea here is immutability. Instead of mutating state in a database record or field, the write is simply recorded, without affecting or destroying any existing information.

Isolation




This has a lot of benefit on backend systems. One is strong isolation between different computations. If one computation has a bug, it doesn't affect the others. And if one computation becomes slow, we can optimize it in isolation.

Swapping computation



We can also swap out different computations. And since we preserve historical information in the log, we can even run computations retroactively, which is useful when building new features or fixing bugs.

More background



<http://blog.confluent.io/2015/03/04/>



<https://twitter.com/jaykrep/status/412643213342568448>



If you're a UI engineer and want more background, these articles are great reads - super clear and accessible.

<http://blog.confluent.io/2015/03/04/>

<https://twitter.com/jaykrep/status/412643213342568448>


The beautiful diagrams here are adapted from Martin Kleppmann's (@martinkl) awesome work.

Applications to UI engineering



So, that's cool, but what does this have to do with UI engineering?

Some similarities...




<http://blog.confluent.io/2015/03/04/>

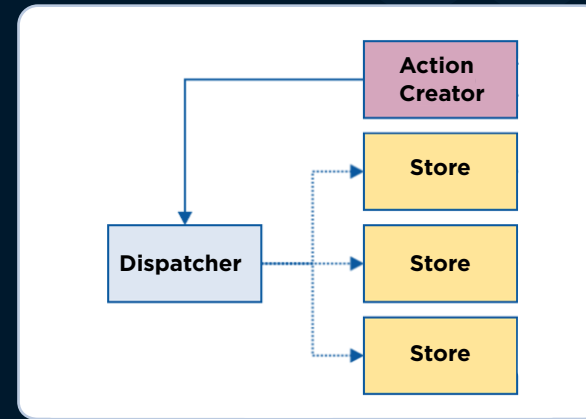


Well, this is the diagram adapted from Samza, a stream processing system.

Some similarities



<http://blog.confluent.io/2015/03/04/>




<http://fluxxor.com/what-is-flux.html>



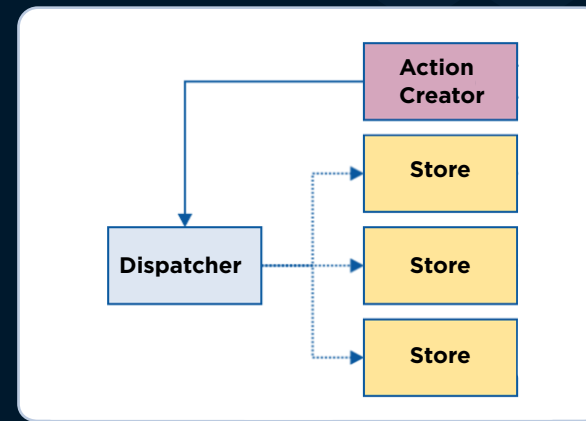
And this diagram is from the documentation for Fluxxor, a great Flux library.

<http://fluxxor.com/what-is-flux.html>

Some similarities



<http://blog.confluent.io/2015/03/04/>




<http://fluxxor.com/what-is-flux.html>




So with Flux, we get similar benefits of isolation between Stores, and the ability to write new Stores without affecting others.

But because Stores are mutable, and because we don't preserve the history of events through the Dispatcher, we can't swap out computation or Stores at runtime.

Some similarities



<http://blog.confluent.io/2015/03/04/>



<http://fluxxor.com/what-is-flux.html>



What if we could? :)

Development superpowers



↓

```
{ "currentSlide": 2, "todos": [] }
```

↓

```
{ "type": "ADD_TODO", "text": "hey" }
```

↓

```
{ "currentSlide": 2, "todos": [ "hey" ] }
```

↓

```
{ "type": "ADD_TODO", "text": "ho" }
```

↓

```
{ "currentSlide": 2, "todos": [ "hey", "ho" ] }
```

Rollback • Commit

<https://github.com/gaearon/redux>



At React Europe, Dan Abramov gave an amazing talk about his work on Redux, which allows hot-reloading of Stores in development mode.

https://twitter.com/dan_abramov/status/617642370591510528?lang=en

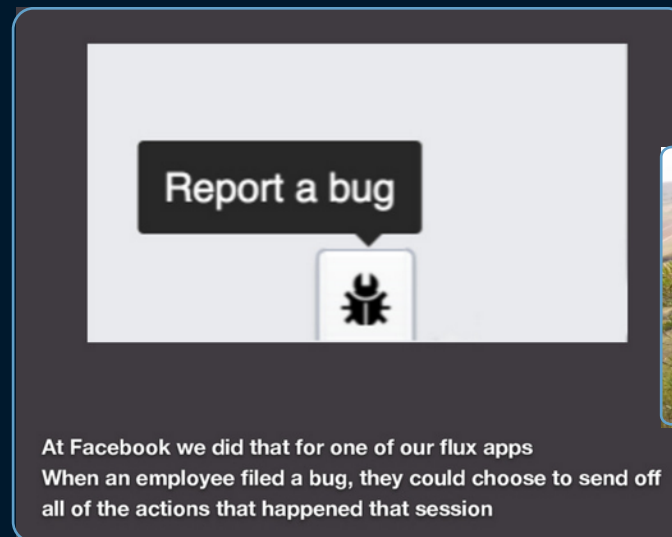
Development superpowers

The screenshot displays the Redux DevTools interface. On the left, a log shows three state snapshots separated by double arrows (⇅). The first snapshot is `{ "currentSlide": 2, "todos": [] }`. The second is `{ "type": "ADD_TODO", "text": "hey" }`. The third is `{ "currentSlide": 2, "todos": ["hey"] }`. Below these is another `{ "type": "ADD_TODO", "text": "ho" }` action, followed by a final state snapshot `{ "currentSlide": 2, "todos": ["hey", "ho"] }`. At the bottom of the log are 'Rollback' and 'Commit' buttons. To the right of the log is a UI mockup titled 'Todos:'. It features an input field, an 'Add todo' button, and a list of items: '• hey' and '• ho'. A small inset video of a speaker is visible in the bottom right corner of the UI mockup area. At the bottom left of the entire image is the URL <https://github.com/gaearon/redux>, and at the bottom right is the Twitter logo.

This works because it's keeping a log of all actions flowing through the dispatcher, and because Stores are expressed as reducers instead of mutable objects.

After hot reloading, redux can re-reduce over the log, which enables hot-swapping Stores.

Debugging superpowers



<https://speakerdeck.com/jmorrell/jsconf-uy-flux-those-who-forget-the-past-dot-dot-dot-1>



Preserving a log helps with debugging as well. Jeremy Morrell has great slides talking about how to take advantage of this in a Flux application.

A separate Store can log all actions flowing through, which lets user sessions be inspected or replayed later.

<https://speakerdeck.com/jmorrell/jsconf-uy-flux-those-who-forget-the-past-dot-dot-dot-1>

Debugging superpowers

Bank Account

-\$10

<https://speakerdeck.com/jmorrell/jsconf-uy-flux-those-who-forget-the-past-dot-dot-dot-1>



Because when debugging, this is what we see using mutable models.

Debugging superpowers

Bank Account

-\$10

Bank Account

Transaction	Amount	Balance
Create Account	\$0	\$0
Deposit	\$200	\$200
Withdrawal	(\$50)	\$150
Deposit	\$100	\$250
		\$250

<https://speakerdeck.com/jmorrell/jsconf-uy-flux-those-who-forget-the-past-dot-dot-dot-1>



And this is what we want to be able to see. The history of what happened and how we got here.

We want to see the log.

Data layer superpowers

```
(aset js/window "undo"  
  (fn [e]  
    (when (> (count @app-history) 1)  
      (swap! app-history pop)  
      (reset! app-state (last @app-history))))))
```

<https://github.com/omcljs/om>
swannodette.github.io/2013/12/31/time-travel/

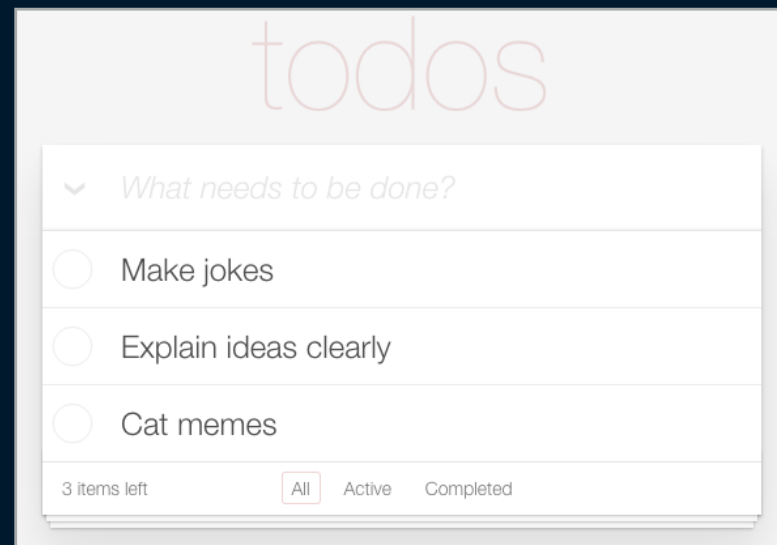


But this can be even more powerful when it comes to building features for end-users like optimistic updates, real-time data systems, or systems for collaboration.

David Nolen has talked about how powerful a feature “undo” can be, and these kinds of features is really where using a log can really pay off.

<http://swannodette.github.io/2013/12/31/time-travel/>

Data layer superpowers



Let's sketch out some examples.

We'll use the universally popular TodoMVC example to see what this might look like.

Embrace the log

timestamp	key	payload
3000	USER_ADDED_TODO	{text: "Make jokes"}
3005	CREATE_TODO_REQUEST	{request: {text: "Make jokes"}}
3500	CREATE_TODO_RESPONSE	{request: {text: "Make jokes"}, response: {id: 14}}
6000	USER_ADDED_TODO	{text: "Explain ideas clearly"}
8000	USER_ADDED_TODO	{text: "Cat memes"}
...		
12000	TODOS_LIST_RESPONSE	[{id: 14, text: "Make funny jokes"}]



Here's an example of a log. As the user is working, their actions are logged, and as requests are made to the server, the requests and responses are logged as well. Whenever the UI "learns" anything at all, it simply records that fact.

This isn't a server-side log being sent down to the UI. It's a log of everything happening in the world, from the UI's perspective.

Embrace the log

timestamp	key	payload
3000	USER_ADDED_TODO	{text: "Make jokes"}
3005	CREATE_TODO_REQUEST	{request: {text: "Make jokes"}}
3500	CREATE_TODO_RESPONSE	{request: {text: "Make jokes"}, response: {id: 14}}
6000	USER_ADDED_TODO	{text: "Explain ideas clearly"}
8000	USER_ADDED_TODO	{text: "Cat memes"}
...		
12000	TODOS_LIST_RESPONSE	[{id: 14, text: "Make funny jokes"}]

TODOS LIST



And the log just sits there. It's immutable, inert.

And we can perform any kind of computation on it that we want, decoupled from when the actions actually occurred.

Embrace the log

timestamp	key	payload
3000	USER_ADDED_TODO	{text: "Make jokes"}
3005	CREATE_TODO_REQUEST	{request: {text: "Make jokes"}}
3500	CREATE_TODO_RESPONSE	{request: {text: "Make jokes"}, response: {id: 14}}
6000	USER_ADDED_TODO	{text: "Explain ideas clearly"}
8000	USER_ADDED_TODO	{text: "Cat memes"}
...		
12000	TODOS_LIST_RESPONSE	[{id: 14, text: "Make funny jokes"}]

TODOS LIST

[{id: 14, text: "Make jokes funny"}]



If we want to simply show the user the list of Todos, we can reduce through the log and pull out the most recent TODOS_LIST_RESPONSE from the server.

This might have come from a specific request made in response to a user action, from polling a REST endpoint, or from data pushed down by a socket.

Embrace the log

timestamp	key	payload
3000	USER_ADDED_TODO	{text: "Make jokes"}
3005	CREATE_TODO_REQUEST	{request: {text: "Make jokes"}}
3500	CREATE_TODO_RESPONSE	{request: {text: "Make jokes"}, response: {id: 14}}
6000	USER_ADDED_TODO	{text: "Explain ideas clearly"}
8000	USER_ADDED_TODO	{text: "Cat memes"}
...		
12000	TODOS_LIST_RESPONSE	[{id: 14, text: "Make funny jokes"}]

TODOS LIST

OPTIMISTIC
TODOS



We can write a function that includes optimistic updates.

Embrace the log

timestamp	key	payload
3000	USER_ADDED_TODO	{text: "Make jokes"}
3005	CREATE_TODO_REQUEST	{request: {text: "Make jokes"}}
3500	CREATE_TODO_RESPONSE	{request: {text: "Make jokes"}, response: {id: 14}}
6000	USER_ADDED_TODO	{text: "Explain ideas clearly"}
8000	USER_ADDED_TODO	{text: "Cat memes"}
...		
12000	TODOS_LIST_RESPONSE	[{id: 14, text: "Make funny jokes"}]

TODOS LIST

OPTIMISTIC
TODOS

```
[ {id: 14, text: "Make jokes funny"},  
  {text: "Explain jokes clearly"},  
  {text: "Cat memes"} ]
```



To do this, we reduce through the log, pulling out both the local user actions and server responses, and merge them both together.

This isn't a super easy function to write, but the point is it's just a function and this is just plain programming.

Embrace the log

timestamp	key	payload
3000	USER_ADDED_TODO	{text: "Make jokes"}
3005	CREATE_TODO_REQUEST	{request: {text: "Make jokes"}}
3500	CREATE_TODO_RESPONSE	{request: {text: "Make jokes"}, response: {id: 14}}
6000	USER_ADDED_TODO	{text: "Explain ideas clearly"}
8000	USER_ADDED_TODO	{text: "Cat memes"}
...		
12000	TODOS_LIST_RESPONSE	[{id: 14, text: "Make funny jokes"}]

TODOS LIST

OPTIMISTIC
TODOS

UNSYNCED
CHANGES



For apps that work offline, we might want to let users see which changes they've made locally, but haven't been broadcast out to the world yet.

Embrace the log

timestamp	key	payload
3000	USER_ADDED_TODO	{text: "Make jokes"}
3005	CREATE_TODO_REQUEST	{request: {text: "Make jokes"}}
3500	CREATE_TODO_RESPONSE	{request: {text: "Make jokes"}, response: {id: 14}}
6000	USER_ADDED_TODO	{text: "Explain ideas clearly"}
8000	USER_ADDED_TODO	{text: "Cat memes"}
...		
12000	TODOS_LIST_RESPONSE	[{id: 14, text: "Make funny jokes"}]

TODOS LIST

OPTIMISTIC
TODOS

UNSYNCED
CHANGES

[{text: "Explain jokes clearly"},
 {text: "Cat memes"}]



Here we just reduce through and collect local changes, and then filter out ones that have already been acknowledged by the server.

Embrace the log

timestamp	key	payload
3000	USER_ADDED_TODO	{text: "Make jokes"}
3005	CREATE_TODO_REQUEST	{request: {text: "Make jokes"}}
3500	CREATE_TODO_RESPONSE	{request: {text: "Make jokes"}, response: {id: 14}}
6000	USER_ADDED_TODO	{text: "Explain ideas clearly"}
8000	USER_ADDED_TODO	{text: "Cat memes"}
...		
12000	TODOS_LIST_RESPONSE	[{id: 14, text: "Make funny jokes"}]

TODOS LIST

OPTIMISTIC
TODOS

UNSYNCED
CHANGES

HIGHLIGHT
CONFLICTS



We can even provide a way for users to see where there have been conflicts between local actions and other users' actions from the server. This doesn't even require special cooperation from the server, even if we use simple polling of a REST API, we can see there's a conflict.

And since we have kept metadata about local timestamps, we have enough information to do more detailed computations. We can pull out the Todos that the user edited in the last ten minutes, and highlight ones where a server update came in that conflicted with the user's change.

Embrace the log

timestamp	key	payload
3000	USER_ADDED_TODO	{text: "Make jokes"}
3005	CREATE_TODO_REQUEST	{request: {text: "Make jokes"}}
3500	CREATE_TODO_RESPONSE	{request: {text: "Make jokes"}, response: {id: 14}}
6000	USER_ADDED_TODO	{text: "Explain ideas clearly"}
8000	USER_ADDED_TODO	{text: "Cat memes"}
...		
12000	TODOS_LIST_RESPONSE	[{id: 14, text: "Make funny jokes"}]

TODOS LIST

OPTIMISTIC
TODOS

UNSYNCED
CHANGES

HIGHLIGHT
CONFLICTS

[{local: {text: "Make jokes"},
server: {id: 14, text: "Make jokes funny"}}]



All of this is possible because we kept the log and full history.

And importantly, this log is from the **user's** perspective. There's an authoritative server truth, but we have enough information in the UI to reconcile changes there with the user's **local perspective** of what has happened.

Conclusion



So that's a sense of this idea. A few more slides, first to acknowledge the challenges here, and then to point you towards some experimentation.

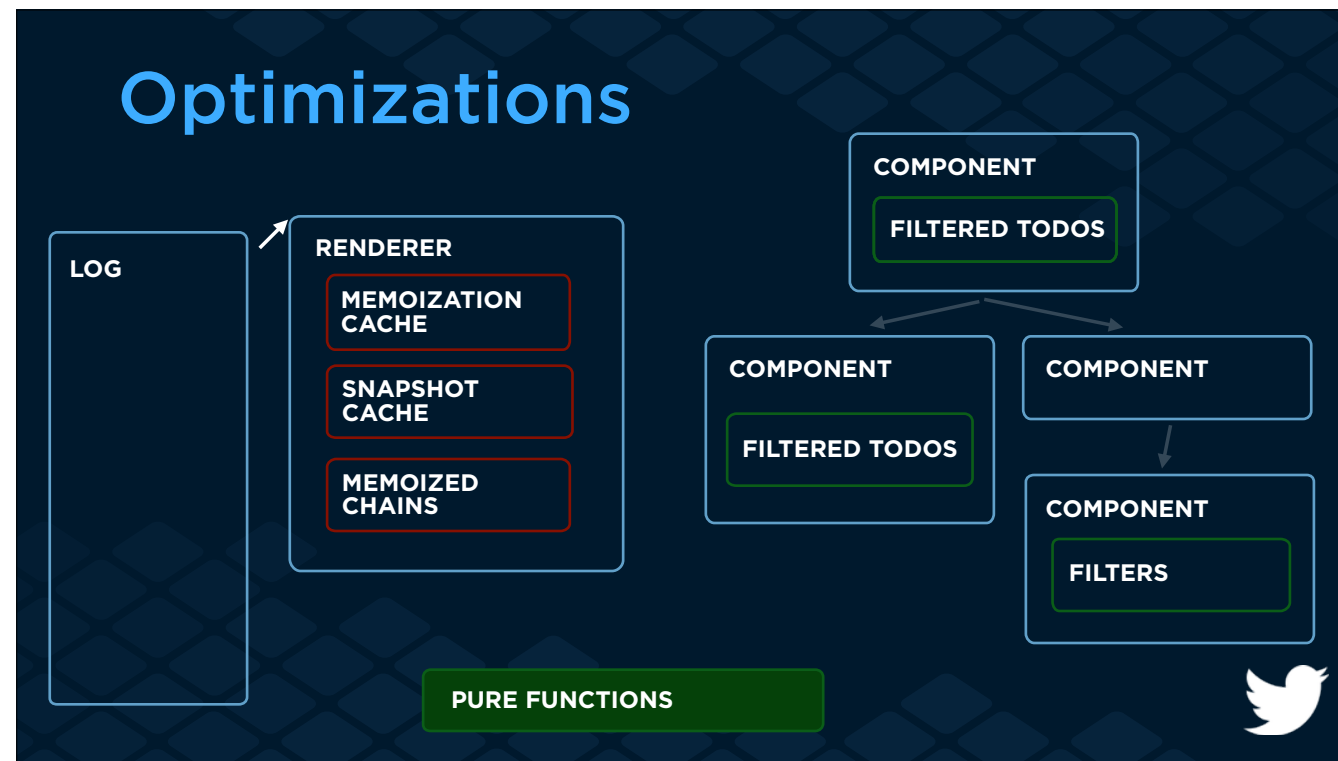
Building blocks



It's one thing to say "well, we can compute anything we want from the log." But it's another to build these abstractions and make them into a developer-friendly API.

This might start with a layer of lower-level computation, but then build up to computing entities, or data for a particular component.

There's work to do here, but this is plain function composition, the easiest kind of programming to refactor, test, re-use, and shape to exactly what you need.



There also might be some performance challenges with a naive solution that keeps a log without bounding its size, or repeatedly reforms a full reduce over the log.

Fortunately, data storage systems in other spaces have tackled these problems, and there's a lot we can learn from there.

See <https://twitter.com/krob/status/617868529082085376> for more information.

Code!

<https://github.com/kevinrobinson/loggit/>

README.md

loggit

Embracing the log for UI engineering!

As data flows through, it's appended to an immutable log. React components describe what computations they want run on that log in a `computations` method. They can then ask the library to perform those computations directly in the render method.

Preserving more information enables features like undo, optimistic updates or surfacing conflicting edits by simply swapping in different computations over the log.



I've put together some code demonstrating these idea, building off Dan's awesome Redux project.

It's on GitHub, and implements some of the optimization strategies mentioned on the previous slide.

<https://github.com/kevinrobinson/loggit/>



If you're curious, there's also some more explanation here: <https://github.com/kevinrobinson/loggit/>.

And my full talk from React Europe is here: https://www.youtube.com/watch?v=EOz4D_714R8

Thanks!

Kevin

@krob